

YAN PAN

San Francisco, CA · yanpan.0508@gmail.com · [GitHub](#)

Professional Experience

Founder & Applied AI Engineer

Oct 2024 – Present

Zen Trading

San Francisco, CA

AI-native quant fund building an agentic quantamental research and trading platform.

- Built an agentic equity-research platform (FastAPI + Next.js) where an LLM agent generates per-company **multi-hop reasoning** theses tracing a company's impact upstream to suppliers and downstream to customers, grounded on point-in-time financials and routed through an observability/audit layer (Langfuse) for traceability
- Productionized **FinRL-X** (the open-source AI4Finance quant framework) around two strategies of our own — an **XGBoost** fundamental stock selector and a **TD3** deep-RL allocator on a shared target-weight interface — adding an Alpaca execution layer and a Dockerized backtest-to-live path; deployed the ensemble to paper trading at a 1.96 Sharpe
- Built **Tradfi-Agent**, an autonomous agent trading tokenized equities on Binance perpetual futures: a two-stage **multi-hop reasoning** loop (macro regime → parallel sector-level LLM calls) emitting long/short decisions with funding-rate-aware entries, a risk gate, and three-layer server-side exit protection over async REST and WebSocket streams

Founding Product Manager

2022 – 2024

FedML / TensorOpera AI

Palo Alto, CA

Product development from research-incubated Federated Learning framework to production Generative AI platform.

- Defined tiered product packaging and GPU pricing for TensorOpera's generative AI platform; partnered with enterprise customers to convert pilots into paid contracts across training and inference workloads
- Shipped core platform features with engineering, streamlining model deployment from a manual multi-hour process into a self-serve workflow; led the AI Agent platform from PRD to GA
- Scaled the Model Hub to cover leading open-weight foundation models (Llama, Mistral, Qwen) with a rapid integration pipeline, broadening the catalog available to enterprise and developer users

AI Product Manager

2021 – 2022

Alibaba Cloud PAI (Platform for AI)

Beijing, China

PAI (Platform for AI) powering Qwen model training, deployment, and inference at scale.

- Pioneered the Federated Learning Privacy module; ran enterprise pilots that converted to paid contracts and secured national-security compliance approvals for regulated-market deployments
- Led launch of the EAS high-availability module — blue-green / canary deployment, auto-scaling, and heterogeneous resource support — targeting a 99.9% uptime SLA for MLOps users
- Designed the EAS pricing model and billing logic, aligning consumption metrics with customer value and supporting platform monetization

Software Engineer Intern

2021

Tencent CV Lab

Shenzhen, China

- Engineered CV inference pipeline for ultra-small instrument defect detection (80→25ms latency, 40 FPS on edge); led MLOps deployment and A/B validation ensuring >95% F1 score and reducing model drift by 40%

Education

Master of Science in Information Technology

2019 – 2022

Carnegie Mellon University

Pittsburgh, PA

Focus on Software Engineering, Data Analytics, and Cloud Computing. Program extended one year due to the U.S. Suspension of Entry as Nonimmigrants of Certain Students and Researchers from the People's Republic of China; worked full-time remotely during the gap year.

Bachelor of Engineering in Intelligent Science and Technology

2015 – 2019

Joint Program of Queen Mary University of London & Beijing University of Posts and Telecommunications

London, UK & Beijing, China

Dual-degree programme jointly delivered by Queen Mary University of London and Beijing University of Posts and Telecommunications — graduates receive one degree from each, taught and assessed in English. Blends UK creative problem-solving with Chinese mathematical rigour. Foundation in computer science, intelligent systems, cryptography, and privacy

Achievements

- **SCALE GMI Startup Cohort** – Accepted into the Apr 2026 cohort
- **NVIDIA Inception Program** – Admitted to NVIDIA's accelerator program for AI-native startups
- **Microsoft Azure Startup Program** – Received \$10,000+ in cloud credits and resources to support AI/ML infrastructure development
- **Lambda Labs Research Grant** – Received GPU compute credits supporting AI/ML research infrastructure
- **Fireworks AI Sponsorship** – Received inference credits for LLM workloads
- **DigitalOcean Sponsorship** – Received cloud infrastructure credits for application hosting and deployment
- **Perplexity AI Pre-IPO Investment** – Participated in pre-IPO funding round as an early investor

Technical Skills

Programming & Data: Python (pandas, numpy, scikit-learn, PyTorch), SQL, API Development (RESTful, gRPC), Git, Docker, CI/CD

AI/ML & Agent Systems: LLM Integration (OpenAI, Anthropic), AI Agent Architectures, Prompt Engineering, RAG Systems, ML Pipeline Development

Product Management: Product Strategy & Roadmap, Stakeholder Management, KPIs & A/B Testing, Cross-functional Leadership, Agile/Scrum

Quantitative Finance: Portfolio Optimization, Options Pricing (Black-Scholes, Greeks), Factor Models, VaR Analysis, Backtesting

Languages

Chinese (Native), English (Fluent)